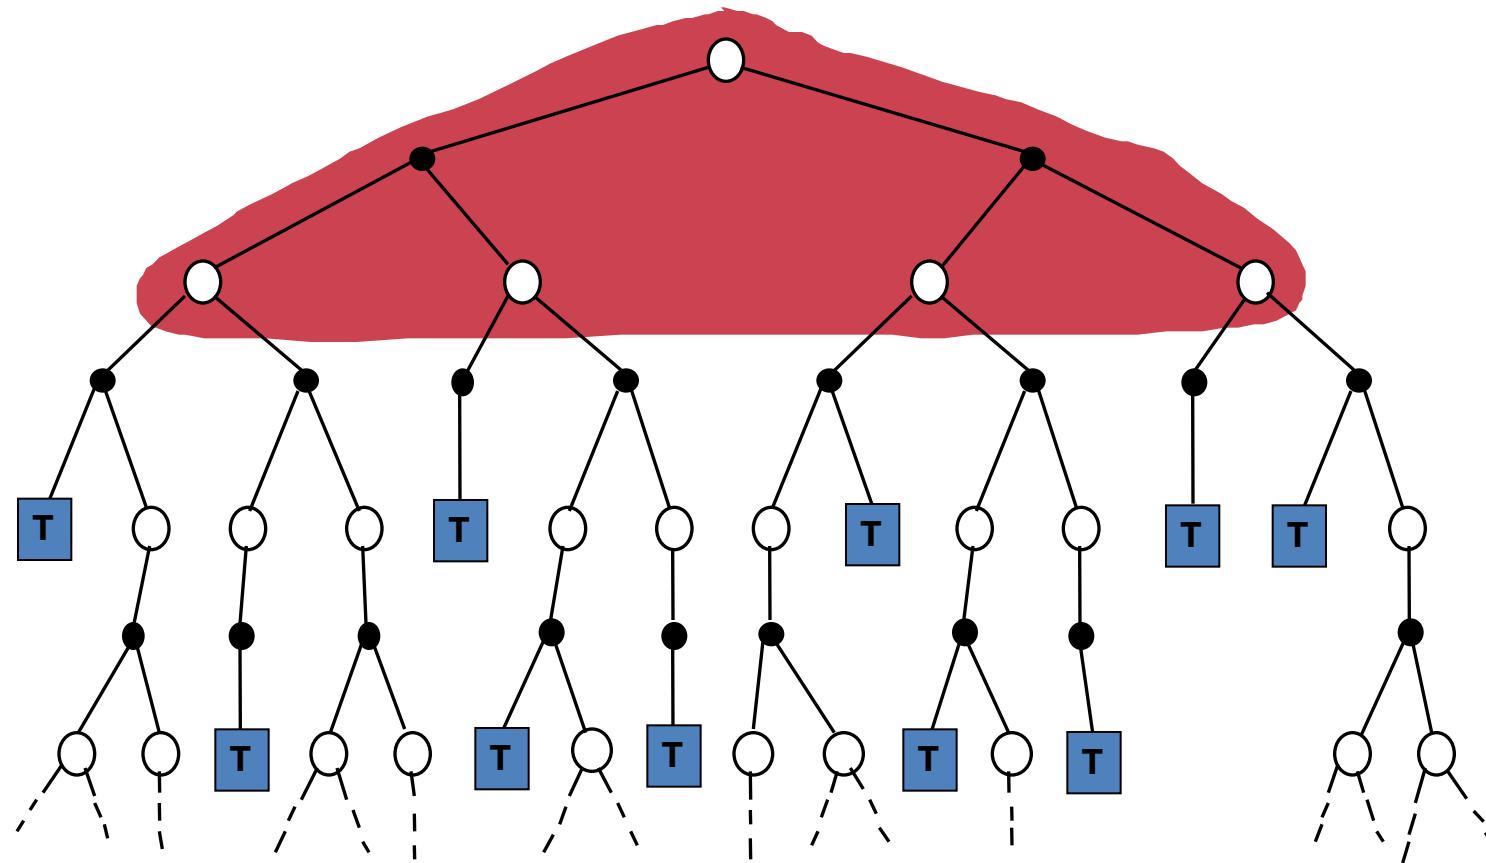


# Lecture 5: Temporal Difference Learning and Monte-Carlo Methods

B. Ravindran

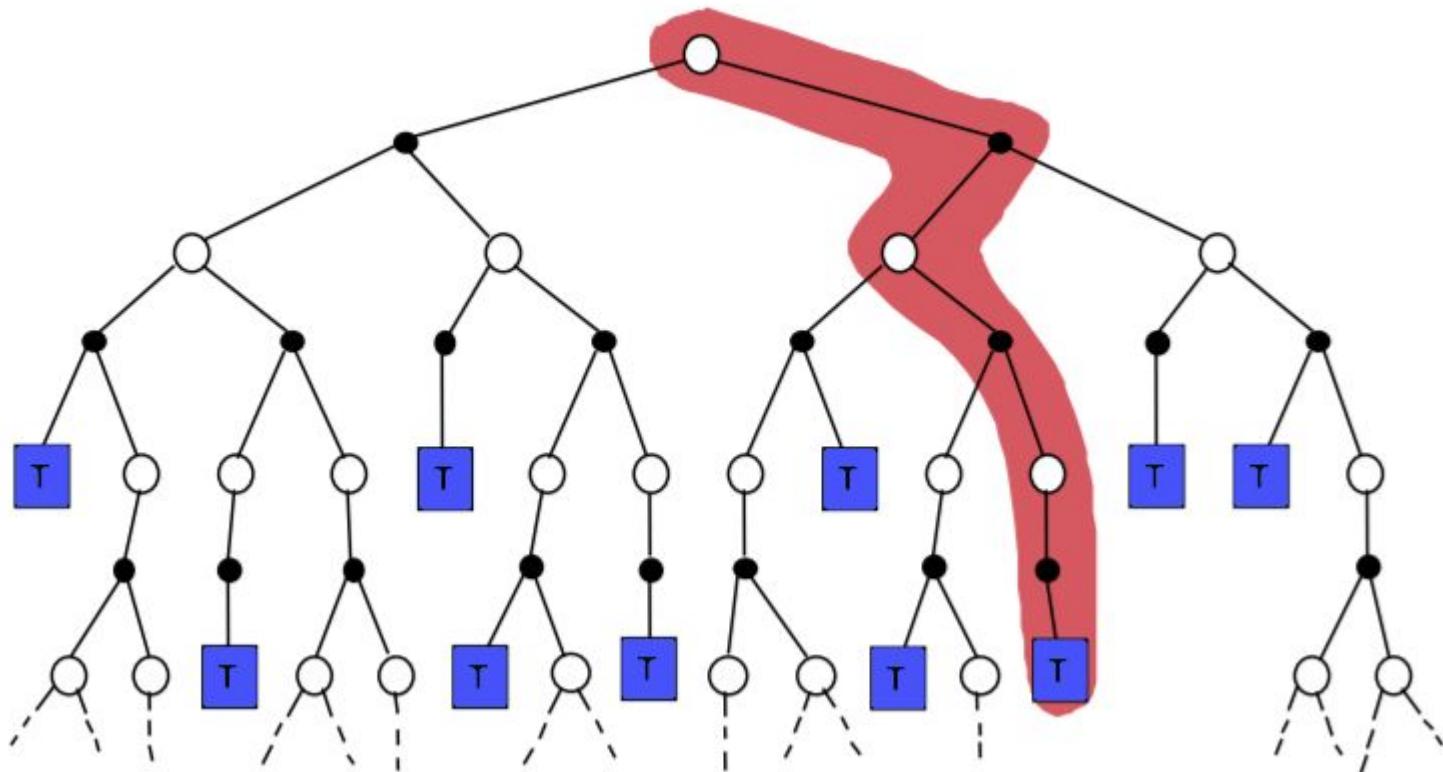
# Dynamic Programming

$$v_{\pi}(s) = \sum_a \pi(a \mid s) \sum_{s',r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

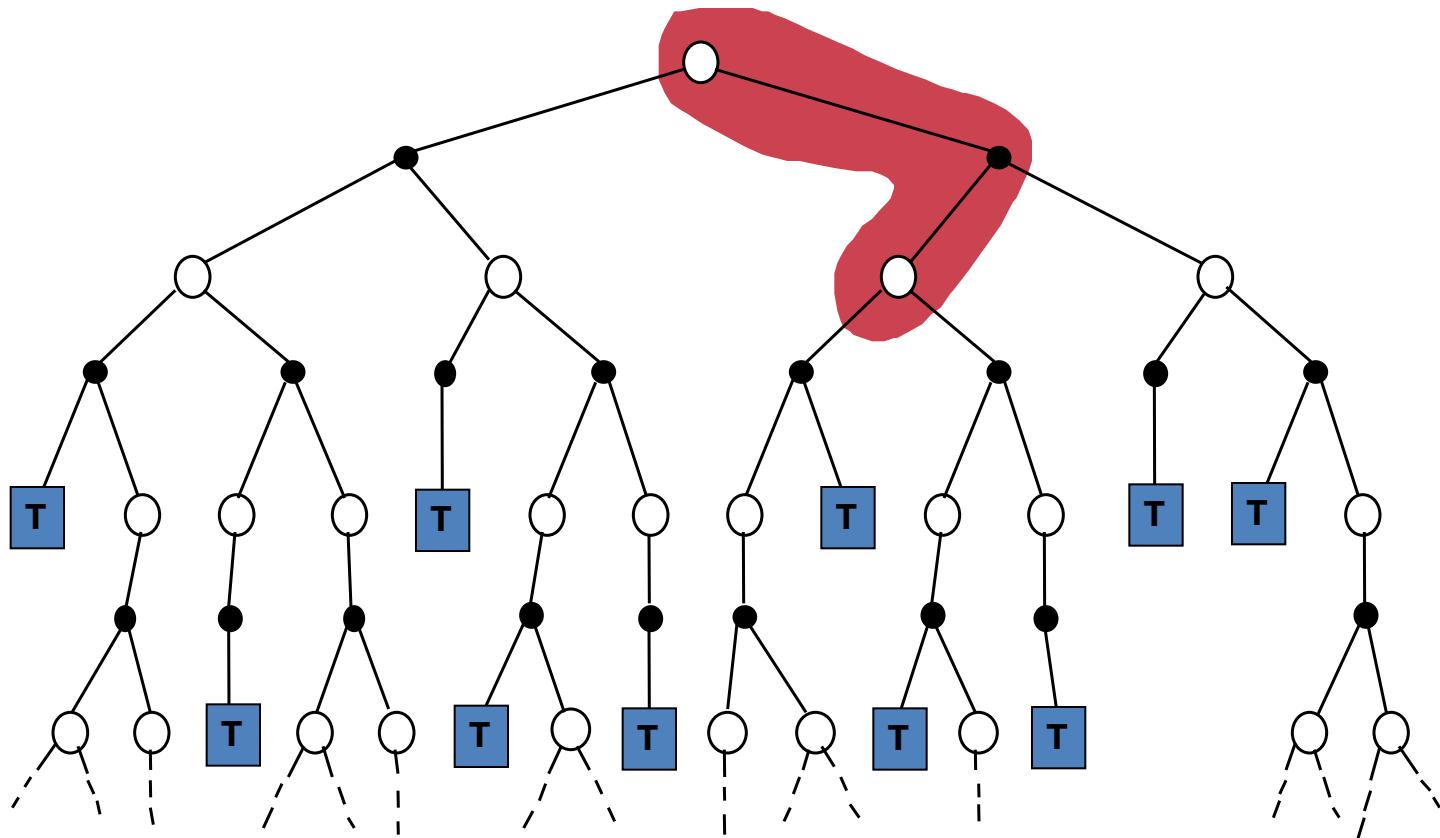


# Monte-Carlo Reinforcement Learning

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s], \text{ for all } s \in \mathcal{S},$$

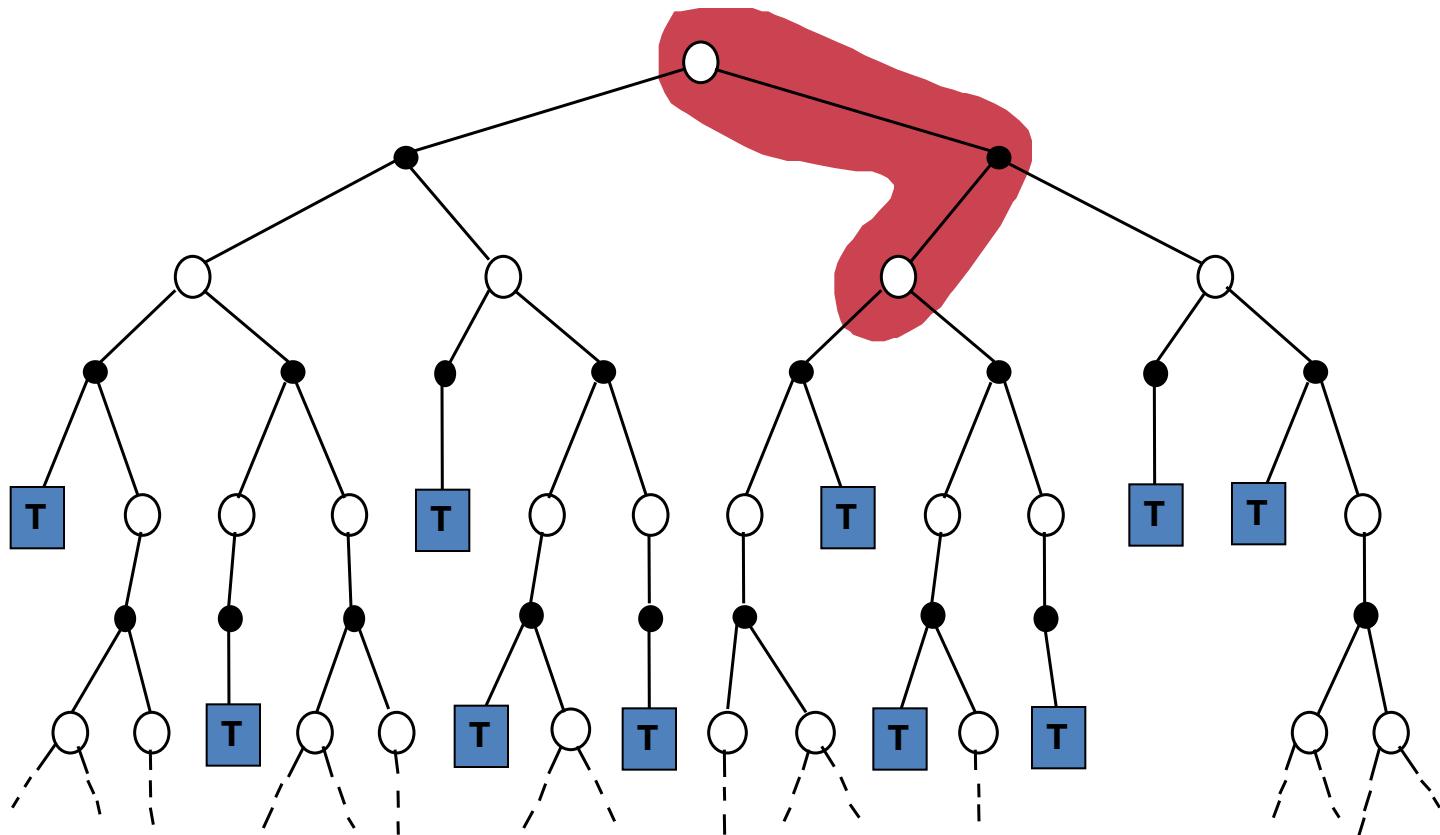


# Simplest “RL” Method



# Simplest “RL” Method

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$



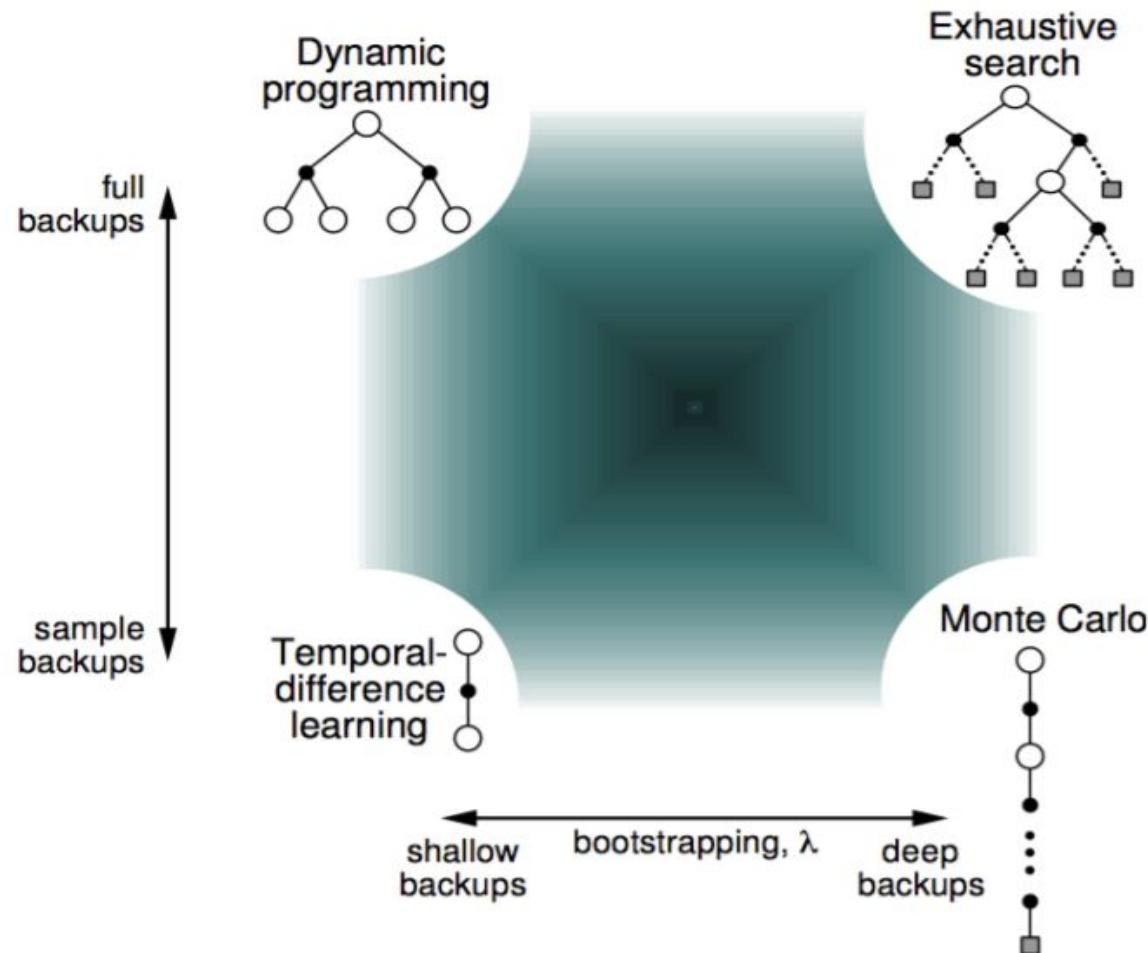
# Temporal Difference

- Simple rule to explain complex behaviors
- Intuition: Prediction of outcome at time  $t+1$  is better than the prediction at time  $t$ . Hence use the later prediction to adjust the earlier prediction.
- Has had profound impact in behavioral psychology and neuroscience!

# Bootstrapping and Sampling

- Bootstrapping: Update using an estimate
  - DP and TD bootstrap
  - Monte Carlo does not bootstrap.
- Sampling: Update calculated using samples without model
  - TD and Monte Carlo sample.
  - DP (typically) does not sample

# Bootstrapping and Sampling



# Monte-Carlo Reinforcement Learning

- Learning directly from sample episodes of experience
- Does not use a known model and is model-free
- MC does not use bootstrapping
- Value functions are calculated as mean of discounted returns  
 $(G_t)$

# Monte-Carlo Prediction

- First-visit MC method estimates  $V(s)$  as the average of the returns following first visits to  $s$ .
- Every-visit MC method averages returns following all visits to  $s$ .

# Monte-Carlo Prediction : First Visit MC

Initialize:

$\pi \leftarrow$  policy to be evaluated

$V \leftarrow$  an arbitrary state-value function

$Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Repeat forever:

Generate an episode using  $\pi$

For each state  $s$  appearing in the episode:

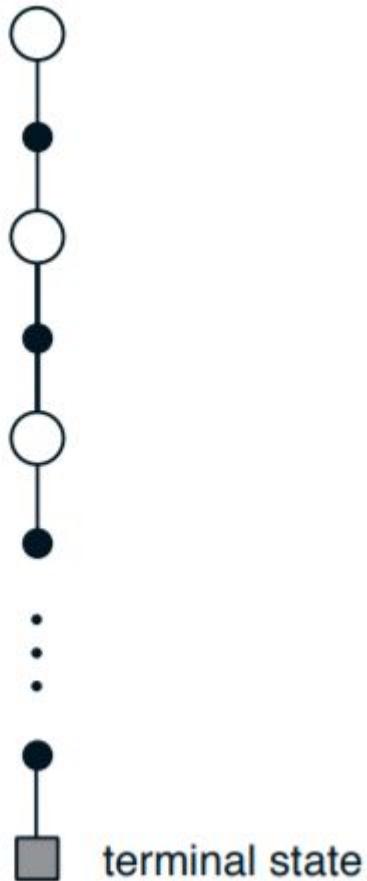
$G \leftarrow$  return following the first occurrence of  $s$

Append  $G$  to  $Returns(s)$

$V(s) \leftarrow$  average( $Returns(s)$ )

---

# Monte-Carlo Prediction : First Visit MC



# Monte-Carlo Control

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow$  arbitrary

$\pi(s) \leftarrow$  arbitrary

$Returns(s, a) \leftarrow$  empty list

Repeat forever:

Choose  $S_0 \in \mathcal{S}$  and  $A_0 \in \mathcal{A}(S_0)$  s.t. all pairs have probability  $> 0$

Generate an episode starting from  $S_0, A_0$ , following  $\pi$

For each pair  $s, a$  appearing in the episode:

$G \leftarrow$  return following the first occurrence of  $s, a$

Append  $G$  to  $Returns(s, a)$

$Q(s, a) \leftarrow$  average( $Returns(s, a)$ )

For each  $s$  in the episode:

$\pi(s) \leftarrow \text{argmax}_a Q(s, a)$

# Bootstrapping and Sampling

- Bootstrapping: Update using an estimate
  - DP and TD bootstrap
  - Monte Carlo does not bootstrap.
- Sampling: Update calculated using samples without model
  - TD and Monte Carlo sample.
  - DP does not sample

# Advantages of TD

- TD methods (like MC) do not require a model of the environment, only experience (sampling)
- TD methods can be fully incremental (bootstrapping)
  - You can learn **before** knowing the final outcome
    - Less memory
    - Less peak computation
  - You can learn **without** the final outcome
    - From incomplete sequences
- TD methods thus combine individual advantages of DP and MC.

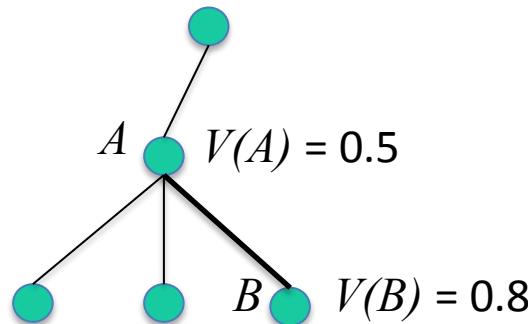
# TD Prediction

- Policy Evaluation (the prediction problem): for a given policy, compute the state-value function.
- No knowledge of  $p$  and  $r$ , but access to the real system, or a “sample” model assumed.
- Uses “bootstrapping” and sampling

The simplest TD method, TD(0):

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

# TD Update Example



Assuming :

reward,  $r$ , A → B : 0

$\alpha$  : 0.2

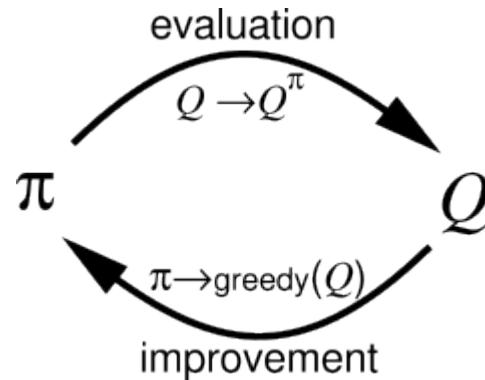
$\gamma$  : 0.9

$$V(A) = V(A) + \alpha[r + \gamma V(B) - V(A)]$$

$$V(A) = 0.5 + 0.2[0 + 0.9 * 0.8 - 0.5] = 0.544$$

# TD Control

- The control problem: approximate optimal policies.
- Recall the idea of GPI:



- Policy evaluation: use TD(0) to evaluate value function.
- Policy improvement: make policy **greedy** wrt current value function.
- Note that we estimate **action values** rather than **state values** in the absence of a model.

# $\varepsilon$ -Greedy Policies

$$a^* \leftarrow \arg \max_a Q(s, a)$$

For all  $a \in \mathcal{A}(s)$ :

$$\pi(a|s) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(s)| & \text{if } a = a^* \\ \varepsilon/|\mathcal{A}(s)| & \text{if } a \neq a^* \end{cases}$$

---

- Any  $\varepsilon$ -greedy policy with respect to  $Q$  following  $\pi$  is an improvement over any  $\varepsilon$ -soft policy  $\pi$  is assured by the policy improvement theorem

# Sarsa: On-Policy TD Control

- In on-policy control, we try improving the policy used for making decisions.

## SARSA

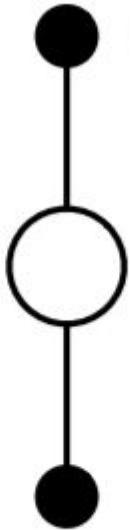
After every transition from a nonterminal state  $s_t$ , do this:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

If  $s_{t+1}$  is terminal, then  $Q(s_{t+1}, a_{t+1}) = 0$ .

- Convergence is guaranteed as long as
  - all state-action pairs are visited an infinite number of times
  - the policy converges in the limit to the greedy policy

# Sarsa: On-Policy TD Control



$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

# Sarsa Algorithm

Initialize  $Q(s,a)$  arbitrarily

Repeat (for each episode)

    Initialize  $s$

    Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\varepsilon$  - greedy)

    Repeat (for each step of episode):

        Take action  $a$ , observe  $r, s'$

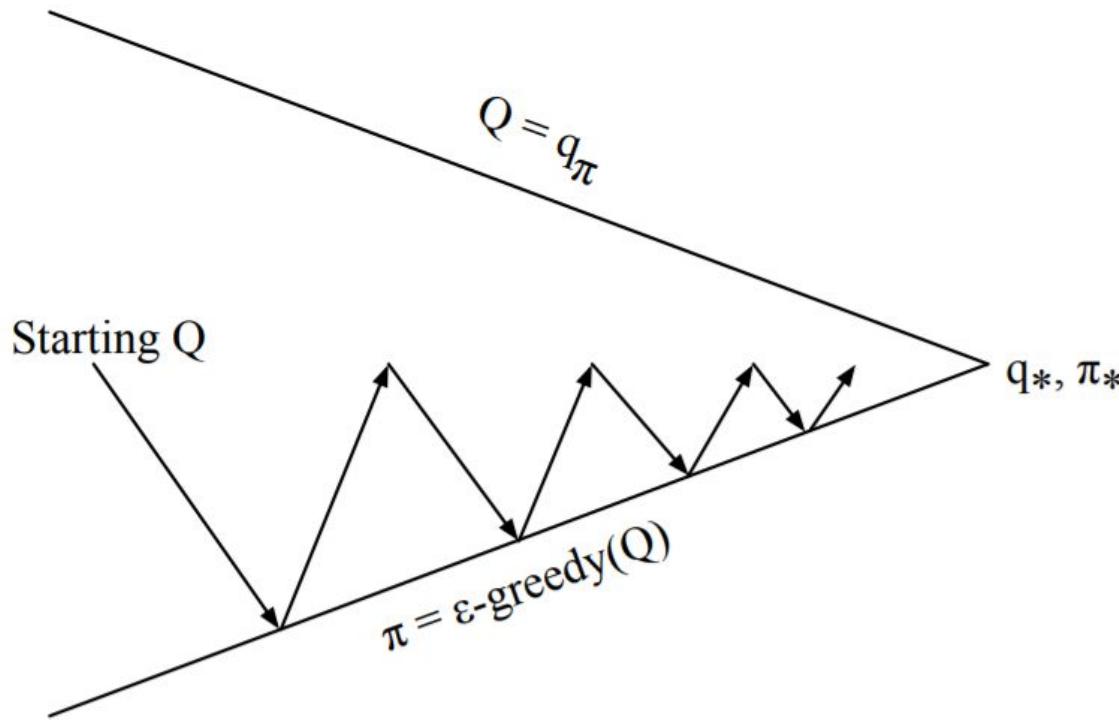
        Choose  $a'$  from  $s'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$  - greedy)

$$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$$

$s \leftarrow s'; a \leftarrow a'$ ;

    until  $s$  is terminal

# Sarsa Algorithm



Every time-step:

Policy evaluation **Sarsa**,  $Q \approx q_\pi$

Policy improvement  $\epsilon$ -greedy policy improvement

# Q-Learning: Off-Policy TD Control

- In off-policy control, we have two policies:
  - the behavior policy – used to generate behavior
  - estimation policy – the policy that is being evaluated and improved.

## Q-learning

One-step Q-learning:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

# Q-learning Algorithm

Initialize  $Q(s, a)$  arbitrarily

Repeat (for each episode) :

    Initialize  $s$

    Repeat (for each step of episode) :

        Choose  $a$  from  $s$  using policy derived from  $Q$  (e.g.,  $\varepsilon$  - greedy)

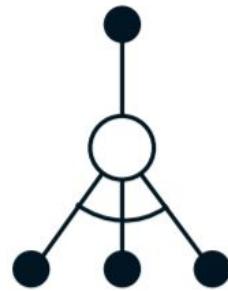
        Take action  $a$ , observe  $r, s'$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

$s \leftarrow s'$ ;

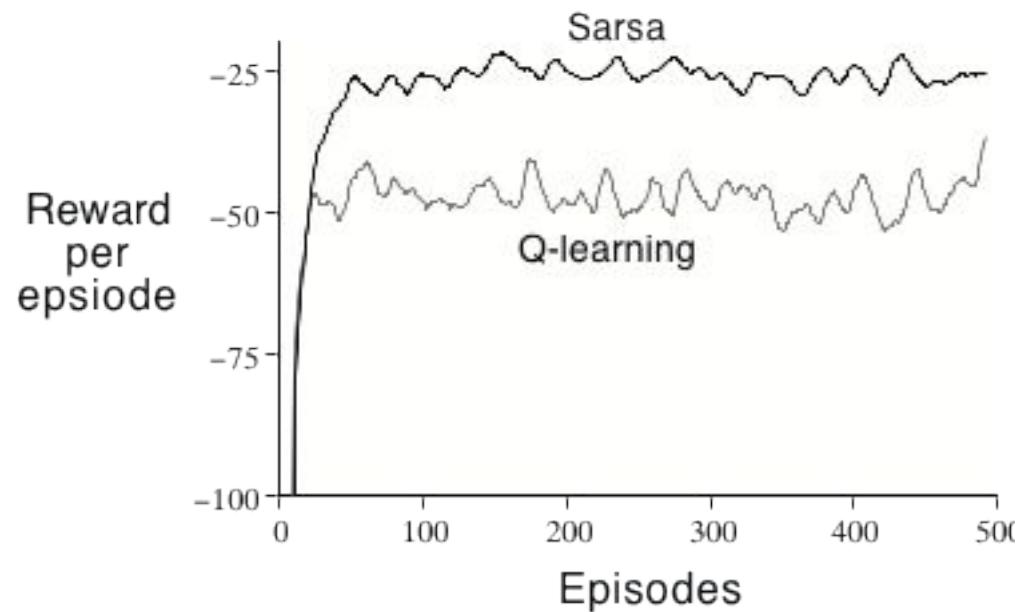
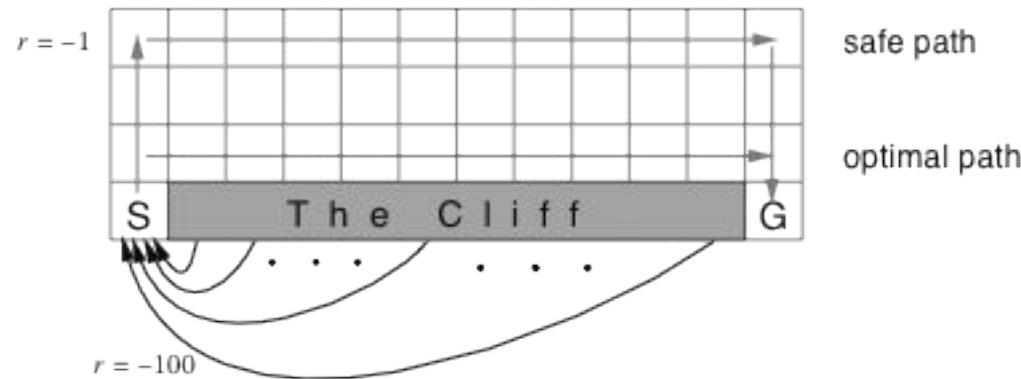
    until  $s$  is terminal

# Q-learning Algorithm



$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

# Cliff Walking: SARSA vs Q-learning



# n-Step TD Prediction

Consider TD(0)

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

Here, the target (for estimating TD error) contains only the next step reward:

$$R_t^{(1)} = r_{t+1} + \gamma V_t(s_{t+1})$$

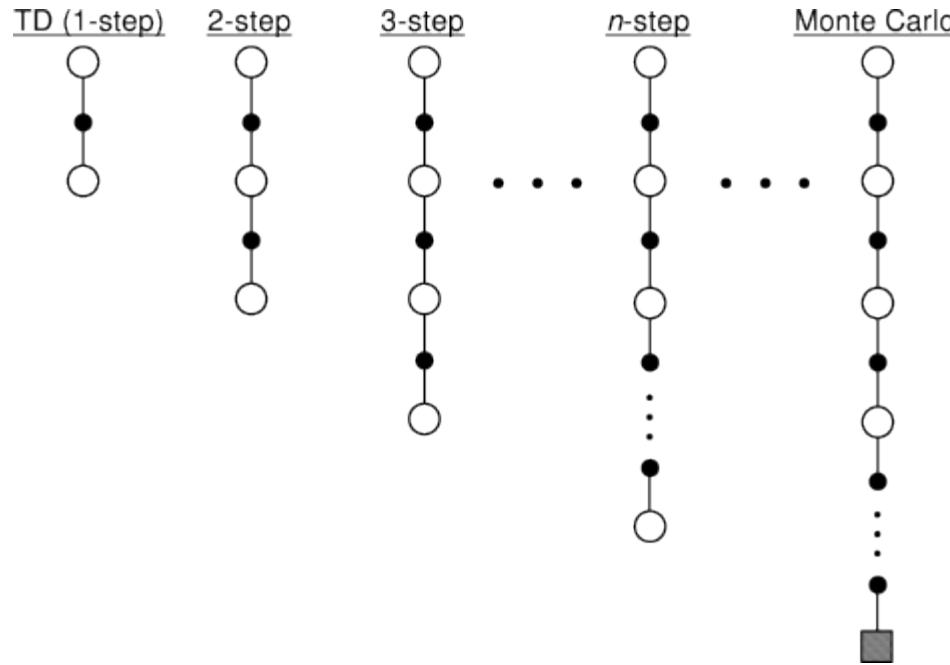
Alternatively, we can consider the rewards received in the next ***n*** steps:

$$V(s_t) \leftarrow V(s_t) + \alpha [r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V(s_{t+n}) - V(s_t)]$$

$$R_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V_t(s_{t+n})$$

The extreme would be to consider rewards till the end of the episode (Monte Carlo).

# $n$ -Step TD Prediction Cont.



The spectrum of back-ups from one-step TD to up-until-termination MC

# TD( $\lambda$ )

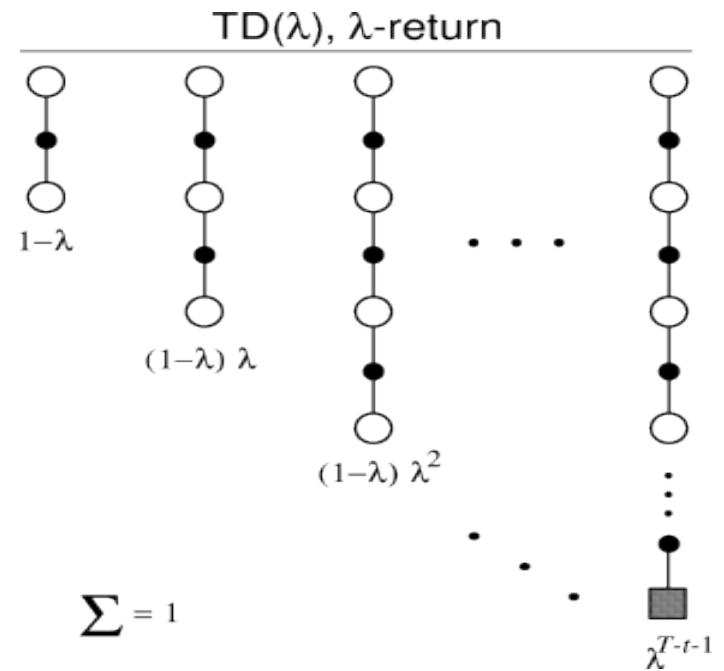
Instead of using  $n$ -step backups, we can consider an average of  $n$ -step returns.

Example:  $R_t^{\text{avg}} = 1/2R_t^{(2)} + 1/2R_t^{(4)}$

In TD( $\lambda$ ), the average contains all the  $n$ -step backups each weighted proportional to  $\lambda^{n-1}$ , where  $0 \leq \lambda \leq 1$ .

$\lambda$ -return:

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_t^{(n)}$$

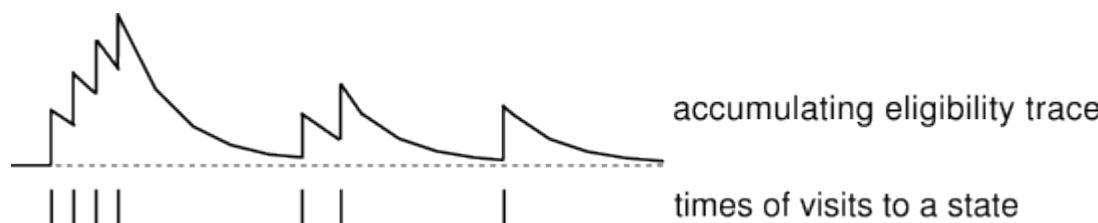


# Eligibility Traces

- To implement TD( $\lambda$ ) we use the concept of eligibility traces.
- These are variables associated with each state denoted by  $e_t(s)$ .
- They indicate the degree to which each state is eligible for undergoing learning changes.
- On each step:

$$e_t(s) = \begin{cases} \gamma \lambda e_{t-1}(s) & \text{if } s \neq s_t; \\ \gamma \lambda e_{t-1}(s) + 1 & \text{if } s = s_t \end{cases}$$

for all  $s \in S$ , where  $\gamma$  is the discount rate.



# TD( $\lambda$ ) Algorithm

Initialize  $V(s)$  arbitrarily and  $e(s) = 0$ , for all  $s \in S$

Repeat (for each episode) :

    Initialize  $s$

    Repeat (for each step of episode) :

$a \leftarrow$  action given by  $\pi$  for  $s$

        Take action  $a$ , observe reward,  $r$ , and next state,  $s'$

$\delta \leftarrow r + \gamma V(s') - V(s)$

$e(s) \leftarrow e(s) + 1$

    For all  $s$  :

$V(s) \leftarrow V(s) + \alpha \delta e(s)$

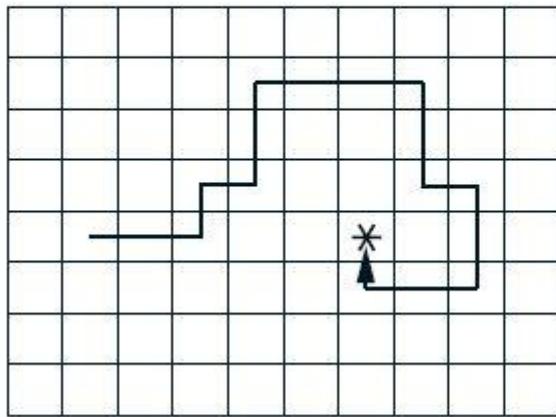
$e(s) \leftarrow \gamma \lambda e(s)$

$s \leftarrow s'$

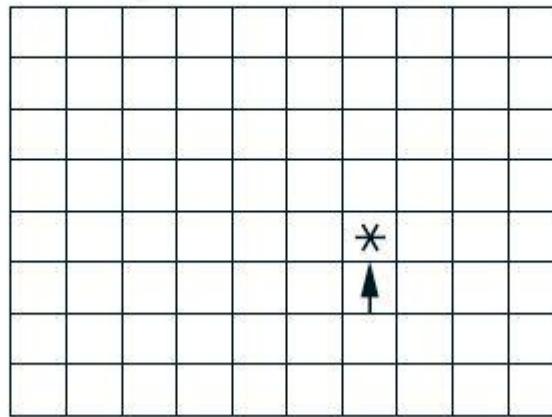
until  $s$  is terminal

# Speedup in Policy Learning

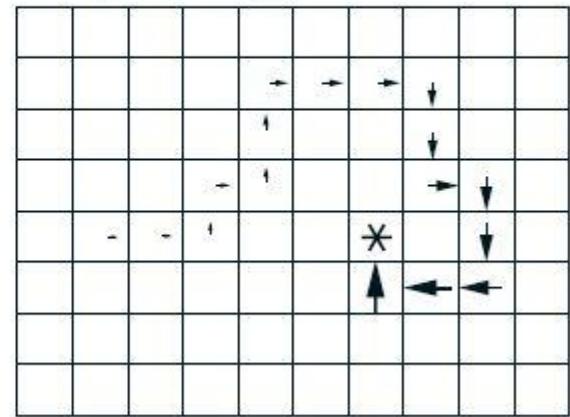
## Path taken



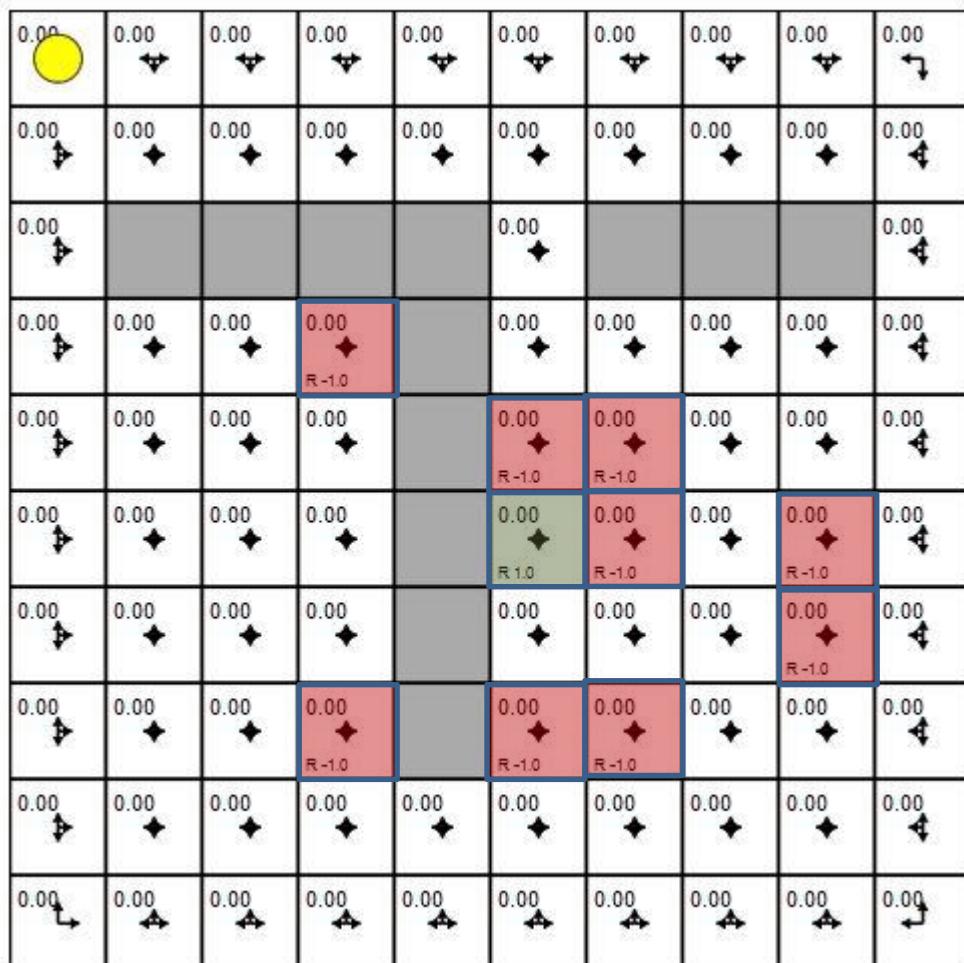
Action values increased by one-step Sarsa



Action values increased by Sarsa( $\lambda$ ) with  $\lambda=0.9$



# TD(0) vs TD( $\lambda$ )



Example grid world domain

Grey cells = obstacles

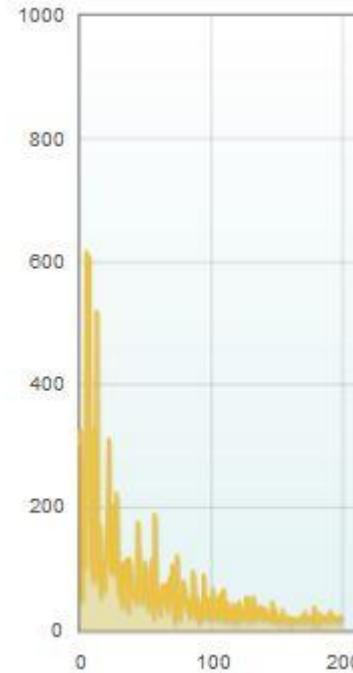
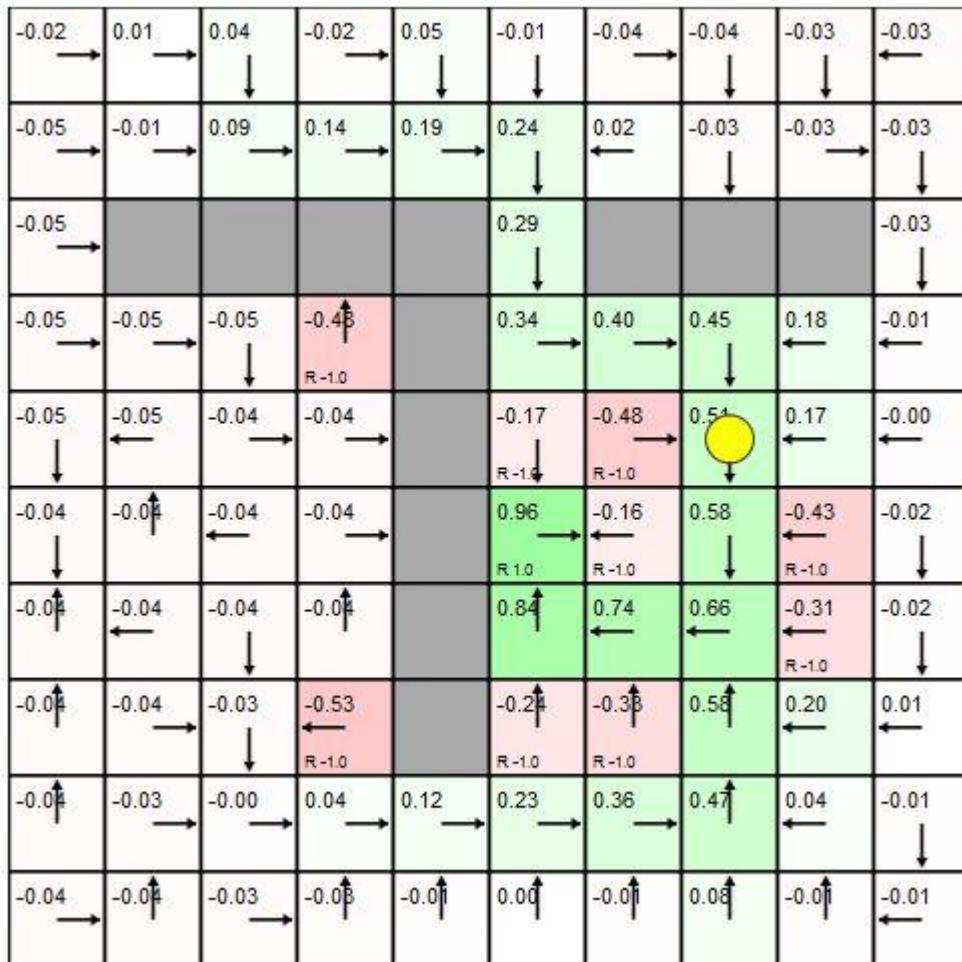
Notice R = 1.0 for centre cell with most surrounding cells having R = -1.0

Cell values = value function estimates

Interactive demo at:

[cs.stanford.edu/people/karpathy/reinforcejs/gridworld\\_td.html](http://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html)

# TD(0) vs TD( $\lambda$ ) Contd.



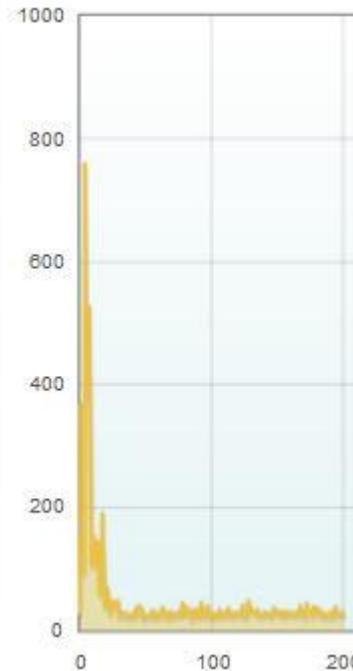
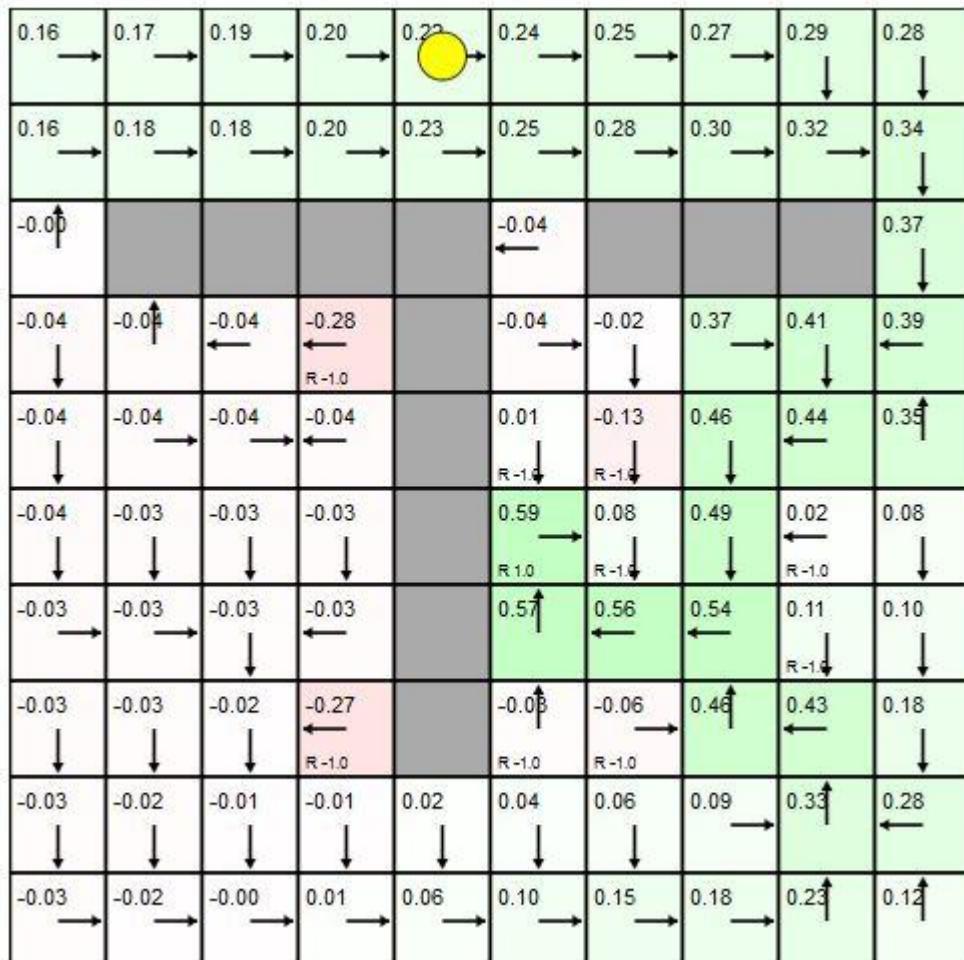
Number of actions before reaching the goal state

Algo : Q - learning

$\gamma = 0.9, \varepsilon = 0.2, \alpha = 0.1, \underline{\lambda = 0}$

Interactive demo at:  
[cs.stanford.edu/people/karpathy/reinforcejs/gridworld\\_td.html](http://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html)

# TD(0) vs TD( $\lambda$ ) Contd.



Number of actions before  
reaching the goal state

Algo : Q - learning

$\gamma = 0.9, \varepsilon = 0.2, \alpha = 0.1, \underline{\lambda = 0.8}$

Interactive demo at:

[cs.stanford.edu/people/karpathy/reinforcejs/gridworld\\_td.html](http://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html)

# More on Off-Policy Learning

- Evaluate target policy  $\pi(a|s)$  to compute  $v_\pi(s)$  or  $q_\pi(s, a)$
- Follow behaviour policy  $\mu(a|s)$
- Assumption of coverage:
  - Every action taken under  $\pi$  is also taken, at least occasionally, under  $\mu$
  - $\pi(a|s) > 0$  implies  $\mu(a|s) > 0$
- e.g.,  $\pi$  can be greedy while  $\mu$  can be  $\varepsilon$ -greedy

# Importance Sampling

- A general technique for estimating expected values under one distribution given samples from another.

$$\begin{aligned}\mathbb{E}_{X \sim P}[f(X)] &= \sum P(X)f(X) \\ &= \sum Q(X) \frac{P(X)}{Q(X)} f(X) \\ &= \mathbb{E}_{X \sim Q} \left[ \frac{P(X)}{Q(X)} f(X) \right]\end{aligned}$$

# Off-Policy MC with Weighted Importance Sampling

Importance-sampling ratio:

$$\rho_t^T = \frac{\prod_{k=t}^{T-1} \pi(A_k|S_k) p(S_{k+1}|S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k|S_k) p(S_{k+1}|S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{\mu(A_k|S_k)}.$$

Weighted average return for  $V$ :

$$V(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)} G_t}{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)}},$$