

Lecture 8

Policy Gradient Algorithms

B. Ravindran

Solution Methods

- Temporal Difference Methods
 - TD(λ)
 - Q-learning
 - SARSA
 - Actor-Critic
- Policy Search
 - Policy Gradient Methods
 - Evolutionary algorithms
- Model based methods
 - Stochastic Dynamic Programming
 - Bayesian approaches

Solution Methods

- Temporal Difference Methods
 - TD(λ)
 - Q-learning
 - SARSA
 - Actor-Critic
- Policy Search
 - Policy Gradient Methods
 - ~~Evolutionary algorithms~~
- Model based methods
 - Stochastic Dynamic Programming
 - ~~Bayesian approaches~~

Solution Methods

- Temporal Difference Methods
 - TD(λ)
 - Q-learning
 - SARSA
 - Actor-Critic
- Policy Search
 - **Policy Gradient Methods**
 - ~~Evolutionary algorithms~~
- Model based methods
 - Stochastic Dynamic Programming
 - ~~Bayesian approaches~~

Policy Search Methods

- Policy search: Instead of maintaining estimates of value functions, search in the space of policies
- **Why?**
 - Simpler description
 - Better convergence
 - Robust to partial observability
- Direct policy search
 - Genetic algorithms
- Policy Gradient Approaches

Policy Gradient Methods

- Policy depends on some parameters θ
 - Action preferences
 - Mean and variance
 - Weights of a neural network
- **Idea:** Modify policy parameters directly instead of estimating the action values
- Maximize:

$$J(\theta) = \mathbb{E}(r_t)$$

Simplified Setting

- θ update:

$$\theta \leftarrow \theta + \alpha \nabla J(\theta)$$

Immediate Reward or
Multi-arm bandits

Multi-Armed Bandit

- n -arm bandit problem is to learn to preferentially select a particular action (arm) from a set of n actions $(1, 2, 3, \dots, n)$
- Each selection results in Rewards, with noise
 - Sometimes you get a reward, sometimes don't
- Pick arm that has the highest expected reward



Policy Gradient Methods

- Policy depends on some parameters θ
 - Action preferences
 - Mean and variance
 - Weights of a neural network
- **Idea:** Modify policy parameters directly instead of estimating the action values
- Maximize:

$$J(\boldsymbol{\theta}) = \mathbb{E}(r_t) = \sum_a Q^*(a)\pi_{\boldsymbol{\theta}}(a)$$

Simplified Setting

Immediate Reward or
Multi-arm bandits

- θ update:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha \nabla J(\boldsymbol{\theta})$$

Stochastic Gradient Descent

- We compute the gradient of the performance $J(\boldsymbol{\theta})$ w.r.t the parameters $\boldsymbol{\theta}$

$$J(\boldsymbol{\theta}) = \mathbb{E} (r_t)$$

$$= \sum_a Q^*(a) \pi_{\boldsymbol{\theta}}(a)$$

$$\nabla J(\boldsymbol{\theta}) = \sum_a Q^*(a) \nabla \pi_{\boldsymbol{\theta}}(a)$$

$$= \sum_a Q^*(a) \frac{\nabla \pi_{\boldsymbol{\theta}}(a)}{\pi_{\boldsymbol{\theta}}(a)} \pi_{\boldsymbol{\theta}}(a)$$

- Estimate the gradient from N samples

$$\hat{\nabla} J(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N r_i \underbrace{\frac{\nabla \pi_{\boldsymbol{\theta}}(a_i)}{\pi_{\boldsymbol{\theta}}(a_i)}}_{\text{Likelihood Ratio}}$$

REINFORCE (Williams '92)

- Incremental version

$$\Delta\theta_t = \alpha r_t \frac{\nabla \pi_{\theta}(a_t)}{\pi_{\theta}(a_t)}$$

$$\Delta\theta_t = \alpha r_t \frac{\partial \ln \pi_{\theta}(a_t)}{\partial \theta}$$

- REINFORCE with baseline

$$\Delta\theta_t = \alpha(r_t - b_t) \frac{\partial \ln \pi_{\theta}(a_t)}{\partial \theta}$$

Reinforcement Baseline Characteristic Eligibility

Policy Gradient Theorem

- In the episodic case, we define the performance by assuming that every episode starts from state s_0 (non-random), as follows:

$$J(\boldsymbol{\theta}) \doteq v_{\pi_{\boldsymbol{\theta}}}(s_0)$$

where $v_{\pi_{\boldsymbol{\theta}}}(s_0)$ is the true value function given a parameterized policy $\pi_{\boldsymbol{\theta}}$

- **The Policy Gradient Theorem:** The gradient of the performance can be expressed in terms of the gradient of the policy, as follows

$$\nabla J(\boldsymbol{\theta}) \propto \sum_s \mu(s) \sum_a q_{\pi}(s, a) \nabla \pi(a \mid s, \boldsymbol{\theta})$$

$\mu(s)$ is the fraction of time spent in state s (on policy distribution)

Policy Gradient Theorem

- The policy gradient theorem provides an analytic expression for the gradient of performance with respect to the policy parameter θ
- We begin the proof by expressing the gradient of the state-value function in terms of the action-value function

$$\begin{aligned}\nabla v_\pi(s) &= \nabla \left[\sum_a \pi(a \mid s) q_\pi(s, a) \right], \quad \text{for all } s \in \mathcal{S} \\ &= \sum_a [\nabla \pi(a \mid s) q_\pi(s, a) + \pi(a \mid s) \nabla q_\pi(s, a)] \\ &= \sum_a \left[\nabla \pi(a \mid s) q_\pi(s, a) + \pi(a \mid s) \nabla \sum_{s', r} p(s', r \mid s, a) (r + v_\pi(s')) \right]\end{aligned}$$

Policy Gradient Theorem: Proof

$$\begin{aligned}\nabla v_\pi(s) &= \sum_a \left[\nabla \pi(a \mid s) q_\pi(s, a) + \pi(a \mid s) \sum_{s'} p(s' \mid s, a) \nabla v_\pi(s') \right] \\ &= \sum_a \left[\nabla \pi(a \mid s) q_\pi(s, a) + \pi(a \mid s) \sum_{s'} p(s' \mid s, a) \right. \\ &\quad \left. \sum_{a'} \left[\nabla \pi(a' \mid s') q_\pi(s', a') + \pi(a' \mid s') \sum_{s''} p(s'' \mid s', a') \nabla v_\pi(s'') \right] \right] \\ &= \sum_{x \in \mathcal{S}} \sum_{k=0}^{\infty} \Pr(s \rightarrow x, k, \pi) \sum_a \nabla \pi(a \mid x) q_\pi(x, a)\end{aligned}$$

$\Pr(s \rightarrow x, k, \pi)$ is the probability of transitioning from state s to state x in k steps under policy π

Policy Gradient Theorem: Proof (contd.)

$$\begin{aligned}\nabla J(\theta) &= \nabla v_\pi(s_0) \\&= \sum_s \left(\sum_{k=0}^{\infty} \Pr(s_0 \rightarrow s, k, \pi) \right) \sum_a \nabla \pi(a \mid s) q_\pi(s, a) \\&= \sum_s \eta(s) \sum_a \nabla \pi(a \mid s) q_\pi(s, a) \\&= \sum_{s'} \eta(s') \sum_s \frac{\eta(s)}{\sum_{s'} \eta(s')} \sum_a \nabla \pi(a \mid s) q_\pi(s, a) \\&= \sum_{s'} \eta(s') \sum_s \mu(s) \sum_a \nabla \pi(a \mid s) q_\pi(s, a) \\&\propto \sum_s \mu(s) \sum_a \nabla \pi(a \mid s) q_\pi(s, a)\end{aligned}$$

REINFORCE: MC Policy Gradient

$$\begin{aligned}\nabla J(\boldsymbol{\theta}) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \boldsymbol{\theta}) \\&= \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a|S_t, \boldsymbol{\theta}) \right] \\&= \mathbb{E}_\pi \left[\sum_a \pi(a|S_t, \boldsymbol{\theta}) q_\pi(S_t, a) \frac{\nabla \pi(a|S_t, \boldsymbol{\theta})}{\pi(a|S_t, \boldsymbol{\theta})} \right] \\&= \mathbb{E}_\pi \left[q_\pi(S_t, A_t) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right] \quad (\text{replacing } a \text{ by the sample } A_t \sim \pi) \\&= \mathbb{E}_\pi \left[G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right], \quad (\text{because } \mathbb{E}_\pi[G_t|S_t, A_t] = q_\pi(S_t, A_t))\end{aligned}$$

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta}_t)}{\pi(A_t|S_t, \boldsymbol{\theta}_t)}$$

Uses the complete return from time t , which includes all future rewards up until the end of the episode

REINFORCE: MC Policy Gradient Algorithm

REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for π_*

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Algorithm parameter: step size $\alpha > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|s, \theta)$

Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$\begin{aligned} G &\leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k && (G_t) \\ \theta &\leftarrow \theta + \alpha \gamma^t G \nabla \ln \pi(A_t | S_t, \theta) \end{aligned}$$

- The algorithm computes an unbiased estimate of the gradient
- Can be very slow due to high variance in the estimates
- Variance is related to the “recurrence time” or the episode length
- For problems with large state spaces, the variance becomes unacceptably high.

REINFORCE: With Baseline

- The policy gradient theorem can be generalized to include a comparison of the action value to an arbitrary baseline $b(s)$

$$\nabla J(\boldsymbol{\theta}) \propto \sum_s \mu(s) \sum_a \left(q_{\pi}(s, a) - b(s) \right) \nabla \pi(a|s, \boldsymbol{\theta})$$

- The baseline should be a function (or a random variable) that does not depend on the action a , in which case, the subtracted quantity is 0

$$\sum_a b(s) \nabla \pi(a|s, \boldsymbol{\theta}) = b(s) \nabla \sum_a \pi(a|s, \boldsymbol{\theta}) = b(s) \nabla 1 = 0.$$

- Update rule of REINFORCE with baseline

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left(G_t - b(S_t) \right) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta}_t)}{\pi(A_t|S_t, \boldsymbol{\theta}_t)}$$

- Doesn't change the expected value, but has an effect on the variance

Actor-Critic Methods

- Actor-Critic methods learn both a policy and a state-value function simultaneously.
- The policy is referred to as the actor that suggests actions given a state.
- The estimated value function is referred to as the critic. It evaluates actions taken by the actor based on the given policy.

$$\begin{aligned}\boldsymbol{\theta}_{t+1} &\doteq \boldsymbol{\theta}_t + \alpha \left(G_{t:t+1} - \hat{v}(S_t, \mathbf{w}) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\&= \boldsymbol{\theta}_t + \alpha \left(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\&= \boldsymbol{\theta}_t + \alpha \delta_t \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)}.\end{aligned}$$

One-step Actor-Critic Algorithm

One-step Actor-Critic (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, w)$

Parameters: step sizes $\alpha^{\theta} > 0$, $\alpha^w > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $w \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Initialize S (first state of episode)

$I \leftarrow 1$

 Loop while S is not terminal (for each time step):

$A \sim \pi(\cdot|S, \theta)$

 Take action A , observe S', R

$\delta \leftarrow R + \gamma \hat{v}(S', w) - \hat{v}(S, w)$ (if S' is terminal, then $\hat{v}(S', w) \doteq 0$)

$w \leftarrow w + \alpha^w \delta \nabla \hat{v}(S, w)$

$\theta \leftarrow \theta + \alpha^{\theta} I \delta \nabla \ln \pi(A|S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$

Monte-Carlo Policy Gradient

- Samples are entire trajectories – $\{s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_T\}$
- Evaluation criterion is the return along the path, instead of the immediate rewards
- The gradient estimation equation hence becomes

$$\hat{\nabla}_\theta \eta(\theta) = \frac{1}{N} \sum_{i=1}^N G_i(s_0) \frac{\nabla_\theta p_{\theta,i}(s_0)}{p_{\theta,i}(s_0)}$$

where, $G_i(s_0)$ is the return starting from state s_0 , and $p_{\theta,i}(s_0)$ is the probability of the i^{th} trajectory, starting from s_0 , and using policy given by θ .

Monte-Carlo Policy Gradient (contd.)

- The “likelihood ratio” in this case evaluates to

$$\frac{\nabla_{\theta} p_{\theta,i}(s_0)}{p_{\theta,i}(s_0)} = \sum_{j=0}^{T-1} \frac{\nabla_{\theta} \pi_{\theta}(s_j, a_j)}{\pi_{\theta}(s_j, a_j)}$$

- The estimate depends on the starting state s_0 . One way to address this issue is to assume a fixed initial state
- More common assumption is to use the average reward formulation

Policy Gradient Theorem

Let $\eta(\Theta) = E_{\pi}(R_t)$

$$\nabla \eta(\Theta) = E_{\pi}[\nabla_{\Theta} \log \pi(s, a; \Theta). Q^{\pi}(s, a)]$$

Can approximate $Q^{\pi}(s, a)$ using the return as we did in MC policy gradient.

Alternatively we can use the Q value itself, learnt via a TD procedure (e.g. SARSA)!

$$\nabla \eta(\Theta) \approx E_{\pi}[\nabla_{\Theta} \log \pi(s, a; \Theta). Q^{\pi}(s, a; w)]$$

$$\Delta \Theta = \alpha \nabla_{\Theta} \log \pi(s, a; \Theta). Q^{\pi}(s, a; w)$$

But the value parameterization should be *compatible* to the action parameterization

Incremental Update

- We can incrementally compute the summation in Equation 1, over one trajectory as follows:

$$z_{t+1} = z_t + \frac{\nabla \pi(s_t, a_t; \Theta)}{\pi(s_t, a_t; \Theta)}$$

$$R_{t+1} = R_t + \frac{1}{t+1} [r_t - R_t]$$

- z_T is known as an eligibility trace. Recall the characteristic eligibility term from REINFORCE:

$$\frac{\partial \ln \pi(a_t; \Theta)}{\partial \Theta}.$$

- z_T keeps track of this eligibility over time, hence is called a trace.

Simple MC Policy Gradient Algorithm

Algorithm 1 Simple MC Policy Gradient Algorithm

- 1: Set $j = 0, R_0 = 0, z_0 = \bar{0}, \Delta_0 = \bar{0}$
- 2: **for** each episode **do**
- 3: **for** each transition s_t, a_t, r_t, s_{t+1} **do**
- 4:
$$z_{t+1} = z_t + \frac{\nabla \pi(s_t, a_t; \Theta)}{\pi(s_t, a_t; \Theta)}$$
- 5:
$$R_{t+1} = R_t + \frac{1}{t+1} [r_t - R_t]$$
- 6: **end for**
- 7: $\Delta_{j+1} = \Delta_j + R_T z_T$
- 8: $j = j + 1$
- 9: **end for**
- 10: **Return** Δ_N/N , where N is the number of episodes

- The algorithm computes an unbiased estimate of the gradient
- Can be very slow due to high variance in the estimates
- Variance is related to the “recurrence time” or the episode length
- For problems with large state spaces, the variance becomes unacceptably high.

Deterministic Policy Gradient (Not part of this lecture)

- Problems with continuous action spaces
 - Hard to implement differentiable continuous controllers in many problems
 - Reduces the expectation to over only the states greatly simplifying the gradient estimation
 - Avoids max over actions
- Key difficulty: What is the notion of gradient for deterministic policies?
 - Continuous action spaces allow us to think of change in actions w.r.t. policy parameter
- Deterministic Policy Gradient Theorem
- <equation>
- Can be shown to be the limit of the Stochastic Policy Gradient Theorem in the limit the variance of the policy goes to zero
- Exploration?
 - If the environment is stochastic then not an issue
 - Otherwise use off-policy actor-critic, where the behavior policy differs from the estimation policy
 - Q – learning for value function updates
- Still require compatible parameterizations