

Actor Critic Methods

B. Ravindran

Recap

- Actor-Critic methods learn both a policy and a state-value function simultaneously.
- The policy is referred to as the actor that suggests actions given a state.
- The estimated value function is referred to as the critic. It evaluates actions taken by the actor based on the given policy.

$$\begin{aligned}\boldsymbol{\theta}_{t+1} &\doteq \boldsymbol{\theta}_t + \alpha \left(G_{t:t+1} - \hat{v}(S_t, \mathbf{w}) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\ &= \boldsymbol{\theta}_t + \alpha \left(R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w}) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)} \\ &= \boldsymbol{\theta}_t + \alpha \delta_t \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)}.\end{aligned}$$

Recap: One Step Actor Critic

One-step Actor–Critic (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, w)$

Parameters: step sizes $\alpha^{\theta} > 0$, $\alpha^w > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $w \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Initialize S (first state of episode)

$I \leftarrow 1$

 Loop while S is not terminal (for each time step):

$A \sim \pi(\cdot|S, \theta)$

 Take action A , observe S', R

$\delta \leftarrow R + \gamma \hat{v}(S', w) - \hat{v}(S, w)$ (if S' is terminal, then $\hat{v}(S', w) \doteq 0$)

$w \leftarrow w + \alpha^w \delta \nabla \hat{v}(S, w)$

$\theta \leftarrow \theta + \alpha^{\theta} I \delta \nabla \ln \pi(A|S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$

Comparison to REINFORCE

- Recall the REINFORCE(with baseline) update:

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left(G_t - b(S_t) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)}$$


- The G_t term (although unbiased) causes the high variance of the algorithm.
- Recall that $\mathbb{E}_\pi[G_t | S_t, A_t] = q_\pi(S_t, A_t)$.
- If we had an estimate of $q_\pi(S_t, A_t)$ with less variance, then we can use that instead of G_t .

Comparison to REINFORCE

- In the one step AC algorithm, we use \hat{v} for both estimating $q_\pi(S_t, A_t)$ and as the baseline.
- The bootstrapping in the update introduces bias but decreases the variance.
- This reduced variance can accelerate learning.

Common Features of AC Methods

- Actor: Computes the policy π_θ and updates θ .
- Critic: Typically computes an estimate $\hat{v}(s, w)$ of the state value function. Updates the parameter w .



Basic Actor Critic Algorithm

1. Take action $\mathbf{a} \sim \pi_\theta(\mathbf{a} \mid \mathbf{s})$ and receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$
2. Update value parameter \mathbf{w} using data $(\mathbf{s}, r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}))$
3. Compute $\hat{\delta}(\mathbf{s}, \mathbf{a}) = r + \gamma \hat{v}(\mathbf{s}', \mathbf{w}) - \hat{v}(\mathbf{s}, \mathbf{w})$
4. $\theta \leftarrow \theta + \alpha \cdot \hat{\delta}(\mathbf{s}, \mathbf{a}) \cdot \nabla_\theta \log \pi_\theta(\mathbf{a} \mid \mathbf{s})$
5. Go back to step 1


How is the critic updated?

- Step 2 of the previous algorithm usually happens in batches. We get multiple data points of the form (\mathbf{s}, y) from parallel workers.
- Minimize the squared loss:


$$L(\mathbf{w}) = \sum_i \|\hat{v}(\mathbf{s}_i, \mathbf{w}) - y_i\|^2$$

Design Choices

two network design



shared network design




Advantage Function

- The advantage function is the difference between the q-value and the value function.
- It can be interpreted as a measure of the advantage of taking action \mathbf{a} in state \mathbf{s} as compared to following policy π

$$\delta_\pi(\mathbf{s}, \mathbf{a}) = q_\pi(\mathbf{s}, \mathbf{a}) - v_\pi(\mathbf{s})$$

A3C – Asynchronous Advantage Actor Critic



A3C - *Mnih et. al. 2016*

Algorithm S3 Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

```

// Assume global shared parameter vectors  $\theta$  and  $w$  and global shared counter  $T = 0$ 
// Assume thread-specific parameter vectors  $\theta'$  and  $w'$ 
Initialize thread step counter  $t \leftarrow 1$ 
repeat
    Reset gradients:  $d\theta \leftarrow 0$  and  $dw \leftarrow 0$ .
    Synchronize thread-specific parameters  $\theta' = \theta$  and  $w' = w$ 
     $t_{start} = t$ 
    Get state  $s_t$ 
    repeat
        Perform  $a_t$  according to policy  $\pi(a_t|s_t; \theta')$ 
        Receive reward  $r_t$  and new state  $s_{t+1}$ 
         $t \leftarrow t + 1$ 
         $T \leftarrow T + 1$ 
    until terminal  $s_t$  or  $t - t_{start} == t_{max}$ 
     $R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, w') & \text{for non-terminal } s_t \end{cases}$  // Bootstrap from last state
    for  $i \in \{t - 1, \dots, t_{start}\}$  do
         $R \leftarrow r_i + \gamma R$ 
        Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; w'))$ 
        Accumulate gradients wrt  $w'$ :  $dw \leftarrow dw + \partial(R - V(s_i; w'))^2 / \partial w'$ 
    end for
    Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $w$  using  $dw$ .
until  $T > T_{max}$ 

```


Reset thread params, update local params with global params

Gather experience

Compute the gradients for this thread



Update global params

A2C – Synchronous Advantage Actor Critic



A3C vs A2C

- We remove the "asynchronous" part of A3C.
- The updates to the global parameters are executed only after all the threads have finished their computation.



Can we re-use experience

- Recall the DQN algorithm where we could use reuse experience using a replay buffer. This was possible as it was an off-policy algorithm.
- Can we do something similar for actor critic algorithms?

Attempt 1

1. Take action $\mathbf{a} \sim \pi_\theta(\mathbf{a} \mid \mathbf{s})$, receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ and store in \mathcal{R}
2. Sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}_{1 \leq i \leq N}$ from buffer \mathcal{R}
3. Update value parameter \mathbf{w} using data $\{(\mathbf{s}_i, r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}))\}_{1 \leq i \leq N}$
4. Compute $\hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) = r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}) - \hat{v}(\mathbf{s}_i, \mathbf{w})$
5. $\theta \leftarrow \theta + \alpha \cdot \left(\frac{1}{N} \sum_{i=1}^N \hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) \cdot \nabla_\theta \log \pi_\theta(\mathbf{a}_i \mid \mathbf{s}_i) \right)$
6. Go back to step 1

Problems?


1. Take action $\mathbf{a} \sim \pi_\theta(\mathbf{a} \mid \mathbf{s})$, receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ and store in \mathcal{R}
2. Sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}_{1 \leq i \leq N}$ from buffer \mathcal{R}
3. Update value parameter \mathbf{w} using data $\{(\mathbf{s}_i, r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}))\}_{1 \leq i \leq N}$
4. Compute $\hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) = r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}) - \hat{v}(\mathbf{s}_i, \mathbf{w})$
5. $\theta \leftarrow \theta + \alpha \cdot \left(\frac{1}{N} \sum_{i=1}^N \hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) \cdot \nabla_\theta \log \pi_\theta(\mathbf{a}_i \mid \mathbf{s}_i) \right)$
6. Go back to step 1

Not the action
taken by π_θ

Not the right target.

Fixes

1. Take action $\mathbf{a} \sim \pi_\theta(\mathbf{a} \mid \mathbf{s})$, receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ and store in \mathcal{R}
2. Sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}_{1 \leq i \leq N}$ from buffer \mathcal{R}
3. Update value parameter \mathbf{w} using data $\{(\mathbf{s}_i, r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}))\}_{1 \leq i \leq N}$
4. Compute $\hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) = r_i + \gamma \hat{v}(\mathbf{s}'_i, \mathbf{w}) - \hat{v}(\mathbf{s}_i, \mathbf{w})$
5. $\theta \leftarrow \theta + \alpha \cdot \left(\frac{1}{N} \sum_{i=1}^N \hat{\delta}(\mathbf{s}_i, \mathbf{a}_i) \cdot \nabla_\theta \log \pi_\theta(\mathbf{a}_i \mid \mathbf{s}_i) \right)$
6. Go back to step 1



Replace \mathbf{a}_i with $\mathbf{a}'_i \sim \pi_\theta(\mathbf{a} \mid \mathbf{s}_i)$

Learn q_π instead. Use the fact that

$$v_\pi(\mathbf{s}) = \mathbb{E}_{\mathbf{a} \sim \pi(\mathbf{a} \mid \mathbf{s})}[q_\pi(\mathbf{s}, \mathbf{a})]$$

Final Algorithm

1. Take action $\mathbf{a} \sim \pi_\theta(\mathbf{a} | \mathbf{s})$, receive $(\mathbf{s}, \mathbf{a}, \mathbf{s}', r)$ and store in \mathcal{R}
2. Sample a batch $\{\mathbf{s}_i, \mathbf{a}_i, r_i, \mathbf{s}'_i\}_{1 \leq i \leq N}$ from buffer \mathcal{R}
3. Update action-value parameter \mathbf{w} using data $\{((\mathbf{s}_i, \mathbf{a}_i), r_i + \gamma \hat{q}(\mathbf{s}'_i, \mathbf{a}'_i, \mathbf{w}))\}_{1 \leq i \leq N}$
4. $\theta \leftarrow \theta + \alpha \cdot \left(\frac{1}{N} \sum_{i=1}^N \hat{q}(\mathbf{s}_i, \mathbf{a}''_i, \mathbf{w}) \cdot \nabla_\theta \log \pi_\theta(\mathbf{a}''_i | \mathbf{s}_i) \right)$
5. Go back to step 1

$$\mathbf{a}''_i \sim \pi_\theta(\mathbf{a} | \mathbf{s}_i)$$

$$\mathbf{a}'_i \sim \pi_\theta(\mathbf{a} | \mathbf{s}_i)$$

Note that we are not subtracting a baseline. This might increase the variance. However this is not an issue as we are averaging the update over a large batch.

Problems?

- The states are not sampled from the steady state distribution of π_θ . This was a requirement for the policy gradient theorem to work.
- Thus, the final policy we get is the optimal policy for a different distribution over the states.
- Similar ideas will be seen when we do soft actor critic.

Compatible Parametrization

- Substituting the approximation $\hat{q}(s, a, w)$ instead of the true value of $q_{\pi}(s, a)$ may introduce bias.
- It can be proved that there is no bias if the function approximator has a “compatible” parametrization with the policy parametrization.
- Condition 1: $\hat{q}(s, a, w) = \nabla_{\theta} \log \pi_{\theta}(a | s)^T w$
- Condition 2: w minimizes mean squared error:

$$w = \arg \min \mathbb{E}_{s \sim \rho^{\pi_{\theta}}, a \sim \pi_{\theta}} \left[(\hat{q}(s, a, w) - q_{\pi_{\theta}}(s, a))^2 \right]$$