



Lecture 12: Advanced Value Based Methods

B. Ravindran

Recall Double Q-learning

Double Q-learning, for estimating $Q_1 \approx Q_2 \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q_1(s, a)$ and $Q_2(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, such that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using the policy ε -greedy in $Q_1 + Q_2$

 Take action A , observe R, S'

 With 0.5 probability:

$$Q_1(S, A) \leftarrow Q_1(S, A) + \alpha \left(R + \gamma Q_2(S', \arg \max_a Q_1(S', a)) - Q_1(S, A) \right)$$

 else:

$$Q_2(S, A) \leftarrow Q_2(S, A) + \alpha \left(R + \gamma Q_1(S', \arg \max_a Q_2(S', a)) - Q_2(S, A) \right)$$

$S \leftarrow S'$

until S is terminal

Different Q functions for selecting an action
and estimating its value



Extending to Double DQN

- Add replay memory and target network
- Make the Q-network a Neural Network
- Add two more networks (as in double Q-learning), one each for the online network and the target network to prevent overestimation????



Extending to Double DQN

- Add replay memory and target network
- Make the Q-network a Neural Network
- Add two more networks (as in double Q-learning), one each for the online network and the target network to prevent overestimation????

We can use the target network to estimate the value while the online network is used for selecting the action!



Extending to Double DQN

Unlike the Q-learning algorithm, where both the networks are updated with equal probability, in Double DQN, only the online network will be updated since the target network is updated towards online network as according to the DQN algorithm.

The equation for the target is therefore:

$$Y_t^{\text{DoubleDQN}} \equiv R_{t+1} + \gamma \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a; \theta_t), \theta_t^-$$

Online network parameters Target network parameters





Prioritized Experience Replay

Idea: Use more important samples more often to converge faster.



Prioritized Experience Replay

Idea: Use more important samples more often to converge faster.

Prob. of sampling transition i

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

Where

- p_i is the priority of transition i
- α determines the prioritisation



Prioritized Experience Replay

Idea: Use more important samples more often to converge faster.

Prob. of sampling transition i

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

Where

- p_i is the priority of transition i
- α determines the prioritisation

Note: $\alpha = 1$ is the default uniform case



Prioritized Experience Replay

Choices for p_i

Case 1: $p_i = |\delta_i| + \epsilon$

Where ϵ is a small positive constant to prevent $p_i=0$

Case 2: $p_i = \frac{1}{\text{rank}(i)}$

Where $\text{rank}(i)$ is the rank of transition i when the replay memory is sorted acc. to $|\delta_i|$.



Prioritized Experience Replay

Problem with $P(i)$:

E(V) expects the updates to come from the same distribution as the samples but PER changes this distribution in an uncontrolled fashion => bias



Prioritized Experience Replay

Problem with $P(i)$:

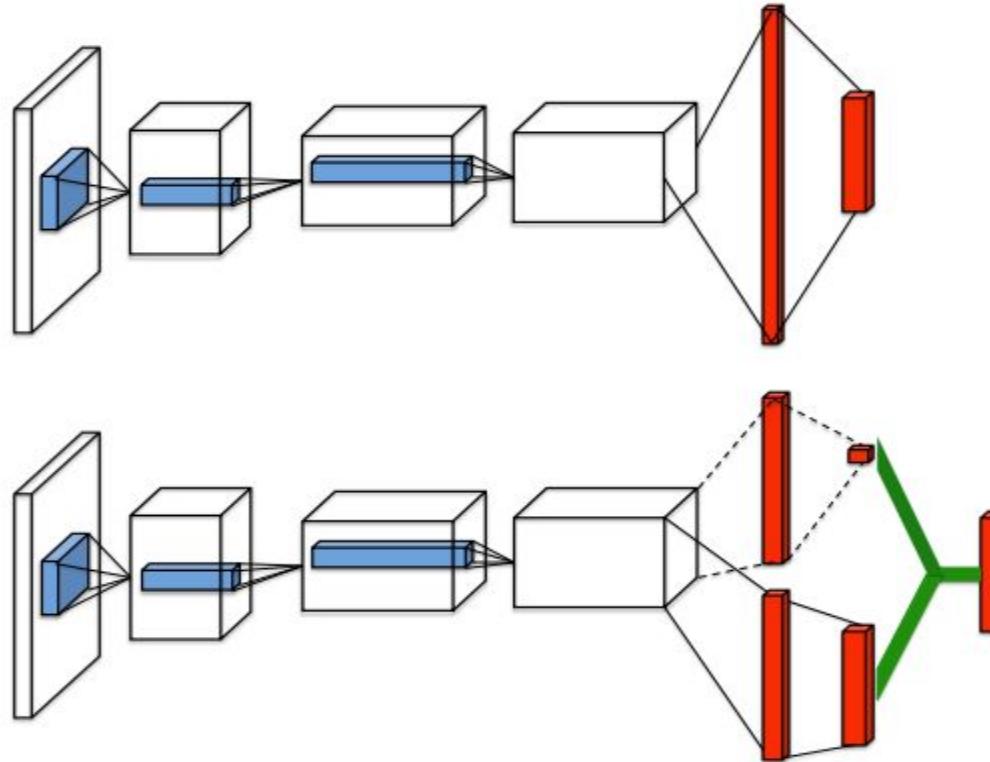
$E(V)$ expects the updates to come from the same distribution as the samples but PER changes this distribution in an uncontrolled fashion \Rightarrow bias

Fixing the bias using *weighted* importance sampling:

$$w_i = \left(\frac{1}{N} \cdot \frac{1}{P(i)} \right)^\beta$$

Use $w_i \delta_i$ instead of δ_i in the Q-learning update

Dueling Network Architecture



Source: [3]

Figure 1. A popular single stream Q -network (**top**) and the dueling Q -network (**bottom**). The dueling network has two streams to separately estimate (scalar) state-value and the advantages for each action; the green output module implements equation (9) to combine them. Both networks output Q -values for each action.



Dueling Network Architecture

The dueling network outputs two streams:

1. Value function (V)
2. Advantage (A)

The aggregation of the two above gives us the Q function, so the network is essentially a Q -network. Thus, we can directly apply Q -learning techniques such as DDQN and PER with the dueling architecture.



Dueling Network Architecture

Aggregation:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + A(s, a; \theta, \alpha)$$

Where α and β are parameters for the two streams of the dueling network.

Does it work?

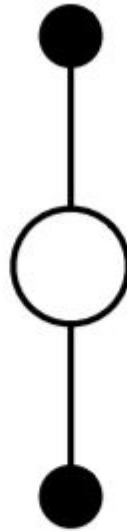


Dueling Network Architecture

Aggregation:

$$Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + \\ \left(A(s, a; \theta, \alpha) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta, \alpha) \right)$$

Recall SARSA



$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$



Recall SARSA

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Problem:

a_{t+1} introduces variance which slows convergence



Expected SARSA

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \mathbb{E}_\pi [Q(S_{t+1}, A_{t+1}) \mid S_{t+1}] - Q(S_t, A_t) \right] \\ &= Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right], \end{aligned}$$

Using the expectation of Q value reduces the variance in the update => we can increase the α to increase the rate of learning.



Expected SARSA

Algorithm 1 Expected Sarsa

- 1: Initialize $Q(s, a)$ arbitrarily for all s, a
- 2: **loop** {over episodes}
- 3: Initialize s
- 4: **repeat** {for each step in the episode}
- 5: choose a from s using policy π derived from Q
- 6: take action a , observe r and s'
- 7: $V_{s'} = \sum_a \pi(s', a) \cdot Q(s', a)$
- 8: $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma V_{s'} - Q(s, a)]$
- 9: $s \leftarrow s'$
- 10: **until** s is terminal
- 11: **end loop**

Source: [4]



References:

- [1] DDQN <https://arxiv.org/pdf/1509.06461.pdf>
- [2] PER <https://arxiv.org/pdf/1511.05952.pdf>
- [3] Dueling <https://arxiv.org/pdf/1511.06581.pdf>
- [4] Expected Sarsa
<http://www.cs.ox.ac.uk/people/shimon.whiteson/pubs/vanseijenadprl09.pdf>