

Soft Actor-Critic





B. Ravindran

Soft Actor-Critic (SAC)

- Stochastic, off-policy, model-free RL algorithm
 - Uses maximum entropy formulation to encourage stability and exploration
 - Sample-efficient
 - Scales to high-dimensional observation/action spaces
 - Robust to random seeds, noise etc.
 - State of the art!
-
- v1 : "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
 - v2 : "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al

Maximum Entropy RL

- Maximize expected return while acting as randomly as possible
- Agent can capture different modes of optimality to improve robustness against environmental changes



Maximum Entropy RL

- Entropy of a random variable x

$$H(P) = \mathbb{E}_{x \sim P} [-\log P(x)].$$

- Maximum Entropy RL objective

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t \left(R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \right]$$

- Here $\alpha > 0$, is the weightage given to the entropy term in the objective. α is commonly referred to as "temperature"

Maximum Entropy RL

- Define the value function to include the entropy bonuses from every timestep:

$$V^\pi(s) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t \left(R(s_t, a_t, s_{t+1}) + \alpha H(\pi(\cdot | s_t)) \right) \middle| s_0 = s \right]$$

- Define the action-value function to include the entropy bonuses from every timestep *except the first*:

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) + \alpha \sum_{t=1}^{\infty} \gamma^t H(\pi(\cdot | s_t)) \middle| s_0 = s, a_0 = a \right]$$

Maximum Entropy RL

- Thus

$$V^\pi(s) = \underset{a \sim \pi}{\mathbb{E}} [Q^\pi(s, a)] + \alpha H(\pi(\cdot | s))$$

- Bellman equation

$$\begin{aligned} Q^\pi(s, a) &= \underset{\substack{s' \sim P \\ a' \sim \pi}}{\mathbb{E}} [R(s, a, s') + \gamma (Q^\pi(s', a') + \alpha H(\pi(\cdot | s')))] \\ &= \underset{s' \sim P}{\mathbb{E}} [R(s, a, s') + \gamma V^\pi(s')] . \end{aligned}$$

SAC

- Policy π_θ
- Two Q functions Q_{w_1} , Q_{w_2}
- Two target Q functions $Q_{w'_1}$, $Q_{w'_2}$
- SAC v1 : Temperature α is a hyperparameter
- SAC v2 : Temperature α is learnt

Learning the Q functions

- Both Q-functions are learned with Mean Squared Bellman Error minimization, by regressing to a single shared target y .

$$L(w_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{D}} [(Q_{w_i}(s,a) - y)^2]$$

- The shared target y is computed using target Q-networks and makes use of the **clipped double-Q** trick.

$$y = r + \gamma \left(\min_{i=1,2} Q_{w'_i}(s', a') - \alpha \log \pi_\theta(a'|s') \right)$$

- The next-state actions used in the target come from the **current policy** instead of the target policy.

$$a' \sim \pi(\cdot|s')$$

Learning the policy

- Maximize

$$\begin{aligned} V^\pi(s) &= \underset{a \sim \pi}{\text{E}} [Q^\pi(s, a)] + \alpha H(\pi(\cdot|s)) \\ &= \underset{a \sim \pi}{\text{E}} [Q^\pi(s, a) - \alpha \log \pi(a|s)]. \end{aligned}$$

- Policy is stochastic, therefore actions are sampled.
- To be able to backprop through sampled actions, we use the **reparameterization trick**
 - Policy outputs mean μ and standard deviation σ of a Gaussian distribution
 - We then sample a gaussian noise $\epsilon \sim \mathcal{N}(0, \mathbb{I})$
 - We combine the noise with the policy outputs and
 - Use tanh to squash the action to [-1,1]

$$a = a_\theta(s, \epsilon) = \tanh(\mu_\theta(s) + \sigma_\theta(s) \cdot \epsilon)$$

- Thus we can rewrite the expectation over actions into an expectation over noise,

$$\mathbb{E}_{a \sim \pi_\theta} [Q^{\pi_\theta}(s, a) - \alpha \log \pi_\theta(a|s)] = \mathbb{E}_{\epsilon \sim \mathcal{N}} [Q^{\pi_\theta}(s, a_\theta(s, \epsilon)) - \alpha \log \pi_\theta(a_\theta(s, \epsilon)|s)]$$

- Thus the final objective becomes

$$\max_{\theta} \mathbb{E}_{\substack{s \sim \mathcal{D} \\ \epsilon \sim \mathcal{N}}} [(\min_{i=1,2} Q_{w_i}(s, a_\theta(s, \epsilon))) - \alpha \log \pi_\theta(a_\theta(s, \epsilon)|s)]$$

Algorithm 1 SAC Algorithm

Initialize networks $\pi_\theta, Q_{w_1}, Q_{w_2}$ with random weights

Initialize target networks $Q_{w'_1} \leftarrow Q_{w_1}, Q_{w'_2} \leftarrow Q_{w_2}$

for Each episode **do**

 Reset environment and get initial state S

for Each time step **do**

 Select action $A \sim \pi_\theta(S)$

 Get next state S' , reward R

 Push (S, A, R, S') into replay buffer \mathcal{D}

 Sample mini batch $\mathcal{B} = \{(s, a, r, s')\}$ from replay buffer.

 Compute target for the Q functions

$$y = r + \gamma \left(\min_{i=1,2} Q_{w'_i}(s', a') - \alpha \log \pi_\theta(a'|s') \right), a' \sim \pi(\cdot|s')$$

 Update Q functions by performing one step of gradient descent on the loss

$$L(w_i) = \frac{1}{|\mathcal{B}|} \sum_{(s,a,r,s') \in \mathcal{B}} [(Q_{w_i}(s, a) - y)^2]$$

 Update policy by performing one step of gradient ascent on the objective

$$J(\theta) = \frac{1}{|\mathcal{B}|} \sum_{s \in \mathcal{B}} [\left(\min_{i=1,2} Q_{w_i}(s, a_\theta(s, \epsilon)) - \alpha \log \pi_\theta(a_\theta(s, \epsilon)|s) \right)]$$

 where $a_\theta(s, \epsilon) = \tanh(\mu_\theta(s) + \sigma_\theta(s) \cdot \epsilon)$ and $\epsilon \sim \mathcal{N}(0, \mathbb{I})$

 Update target networks

$$w'_1 \leftarrow \tau w_1 + (1 - \tau) w'_1$$

$$w'_2 \leftarrow \tau w_2 + (1 - \tau) w'_2$$

$$S \leftarrow S'$$


end for

end for


SAC v1

- "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor", Haarnoja et al
- Temperature α is a hyperparameter
- Tasks
 - A range of continuous control tasks from the OpenAI gym benchmark suite
 - RL-Lab implementation of the Humanoid task
 - The easier tasks can be solved by a wide range of different algorithms. The more complex benchmarks, such as the 21-dimensional Humanoid (rllab) are exceptionally difficult to solve with off-policy algorithms.
- Baselines:
 - DDPG, SQL, PPO, TD3 (concurrent)
 - TD3 is an extension to DDPG that first applied the double Q-learning trick to continuous control along with other improvements.


SAC v1 : Experimental Results




(a) Hopper-v1




(b) Walker2d-v1




(c) HalfCheetah-v1



(d) Ant-v1




(e) Humanoid-v1




(f) Humanoid (rllab)

SAC v1 : Ablation Study


- How does the stochasticity of the policy and entropy maximization affect the performance?
- Comparison with a deterministic variant of SAC that does not maximize the entropy and that closely resembles DDPG




SAC v1 : Hyperparameter Sensitivity



(a) Evaluation



(b) Reward Scale



(c) Target Smoothing Coefficient (τ)

Limitation of SAC v1

- SAC v1 is brittle to the choice of temperature α , a hyperparameter that controls exploration
 - Solution --> Automatic temperature tuning!

SAC v2

- "Soft Actor-Critic: Algorithms and Applications", Haarnoja et al
- Temperature α is learnt by minimizing the loss

$$L(\alpha) = \alpha (-\log \pi(a|s) - \tilde{H})$$

where \tilde{H} is the entropy target.

- Typically, \tilde{H} it is set to be equal to the negative of the action space dimension i. e. $\tilde{H} = -\dim(\mathcal{A})$
- SAC v2 shows results on simulated tasks from OpenAI gym, RL Lab as well as real-world tasks such as locomotion for a quadrupedal robot and robotic manipulation with a dexterous hand

SAC v2 : Experimental Results - RL Lab





Figure 1: Training curves on continuous control benchmarks. Soft actor-critic (blue and yellow) performs consistently across all tasks and outperforming both on-policy and off-policy methods in the most challenging tasks.

SAC v2 : Real World Robots

Quadrupedal Robot Locomotion




SAC v2 : Real World Robots

- Dexterous hand manipulation
 - 20 hour end-to-end learning
 - Valve position as input: SAC 3 hours vs. PPO 7.4 hours



Limitations/Open Issues

- Lack of experiments on hard-exploration problems
- High-variance due to automatic temperature tuning



Recap: SAC

- A stochastic, off-policy, model-free maximum entropy deep RL algorithm
 - Sample-efficient
 - Scales to high-dimensional observation/action spaces
 - Robust to random seeds, noise etc.
- SAC outperforms SOTA model-free deep RL methods, including DDPG, PPO in terms of average return, sample complexity and robustness.

Thank you !!