


Introduction to Hierarchical Reinforcement Learning


B. Ravindran

- Why function approximation?
 - ‘Scale up’ the solution methods for solving large problems
 - ‘Trying’ to find structure in the problem (Generalization)
- Another approach:
 - We can look for a ‘higher order structure’.
[Sub-problems that repeatedly occurs as part of the larger problem
Example: Navigation]



TEMPORAL ABSTRACTION:


- These new sub-problems are obtained by abstracting away the notion of time.



- Solving 'Sub-problems' (+) that could take 'varying amounts of time'.

- **TRANSFER/REUSABILITY:**


- These ‘sub-problems’ can be broken down into even smaller problems.



- These ‘Sub-problems’ tend to re-occur very often.
(as sub-parts of other problems as well)
- Hierarchies could give us ‘Meaningful Modules’ to optimize.



- **MEANINGFUL ABSTRACTIONS:**

- Optimal solution for a ‘sub-problem’ does not depend on other parallel sub-problems.





- This way, we have a simpler problem to solve.
(State-space of a sub problem is way smaller than a bigger problem)
- Efficient learning of sub problems (+) Reusability for other problems.

Optimality



Which of the following is the 'optimal' policy for an agent starting from Room A that wishes to reach 'G'?

Optimality




Which of the following is the 'optimal' policy for an agent starting from Room A that wishes to reach 'G'?


Both 'LEFT' & 'RIGHT' could be the answer depending on the which optimality is considered.

- Hierarchically Optimal & Recursively Optimal

Recursively Optimal (module-wise best way)



Hierarchically Optimal (Overall best way)




- Search through policies that respect the hierarchy. (Leave Room A to go to next room)
- Hierarchically Optimal policies are policies that gives the best solution yet satisfies the proposed hierarchy
- Recursively Optimal policies are the best policies formed by putting together component-wise policies obtained by optimizing for each component.
- Flat Optimality -> Optimal policy (which does not consider any hierarchy)

Goal: A \rightarrow C

Hierarchy presented:

$$A \rightarrow C : (A \rightarrow B) + (B \rightarrow C)$$

Recursively Optimal
Hierarchically Optimal
Flat Optimal




Comparison

- Cost-wise:
 - Flat (better than) Hierarchical (better than) Recursive
- Flat optimal is difficult to obtain.
- Recursive is better than hierarchical?
 - Modular solutions (reusable)
 - Learning horizon could be simpler (if problem can be broken down into meaningful subproblems)
 - The optimal sub-solutions put together could be bad.
- Recursive Optimality works well if there is a mechanism that can treat the ‘sub-problem’ solutions as actions that can be used to reduce the learning horizon of the bigger problem. SMDPs help us here.


SMDP - (Semi Markov Decision Process)

- Actions have durations (τ - holding time)

MDP



SMDP



- Each action could have different time duration. (τ & τ')

SMDPs

Transitions: $\langle S, A, P, R \rangle$

One-step transition probabilities:

$$p(s', \tau | s, a) \quad \vdash \quad \begin{array}{l} p(s' | s, a) \\ p(\tau | s, a) \end{array}$$

One-step expected-rewards:

$$E[r | s, a, s', \tau]$$

SMDP Q-Learning

- Q-Update Equation (SMDPs):

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha \left[\bar{r}_{t+\tau} + \gamma^\tau \max_{a'} Q(s_{t+\tau}, a') - Q(s_t, a_t) \right]$$

- The ‘reward’ used in the above update (reward obtained from the update sample with ‘tou’ timestep) can be framed as

$$\bar{r}_{t+\tau} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{\tau-1} r_{t+\tau}$$