

HRL: Options

B. Ravindran

Options Framework



- Macro Actions: Sequence of actions put together to be used as a ‘single’ action.
- Options:
 - Initiation Set $\mathcal{I}_o \subseteq S$
 - Policy π_o
 - Termination $\beta_o : S \rightarrow [0, 1]$

Types of Options

- Markov Options:
 - π_o depends only on current state
- Semi-Markov Options:
 - π_o depends on history since option started

example:

Semi-Markov Option
 $\pi_o: (\text{Right}) + (\text{Right}) + (\text{Up})$



Learning with Options

- SMDP Q-Learning:
 - In SMDP Q-Learning, if a primitive action was selected in a state, the value of the state-action pair is updated according to the regular Q-Learning update rule.
 - If the agent selected an option o , no state-action values are updated until o terminates.
 - At this point, the cumulative, discounted reward received during the execution of the option is used to update the value of the option in the state s in which it was initiated.

$$Q(s_t, a_t) = Q(s_t, a_t) + \alpha \left[\bar{r}_{t+\tau} + \gamma^\tau \max_{a'} Q(s_{t+\tau}, a') - Q(s_t, a_t) \right]$$

$$\bar{r}_{t+\tau} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{\tau-1} r_{t+\tau}$$

Intra-Option Q-Learning

- At every step, the state-action value for the primitive action as well as the state-action value for all options that would have selected the same action are updated, regardless of the option in effect.
- Let $\pi_o = a_1(\text{for } s_1), a_2(\text{for } s_2) \dots \dots$
- For options,

$$\begin{aligned} Q(s_1, o) &= Q(s_1, o) + \alpha[r_1 + \gamma Q(s_2, o) - Q(s_1, o)] && \text{If not terminating at } s_2 \\ &= Q(s_1, o) + \alpha \left[r_1 + \gamma \max_a Q(s_2, a) - Q(s_1, o) \right] && \text{If terminating at } s_2 \end{aligned}$$

[Tutorial] - Implement as a single function with termination probability

Learning with Options

- For primitive actions (state-action pairs), we use regular Q-learning update.

$$Q(s_1, a_1) = Q(s_1, a_1) + \alpha \left[r_1 + \gamma \max_a Q(s_1, a) - Q(s_1, a_1) \right]$$
$$Q(s_2, a_2) = \dots$$

- Additionally, an option execution allows us to update for all other options that are consistent with the first option. (for every other option o' that would have selected the same action a in particular states)
- Suppose $\pi_o(s_1) = a^\# \& \pi_{o'}(s_1) = a^\#$
- Even when executing option o , o' can be updated


$$Q(s_1, o') = Q(s_1, o') + \alpha(r_1 + \gamma Q(s_2, o') - Q(s_1, o'))$$

Option Discovery

- What makes an option ‘good’?
 - Reusability
 - Cuts down on exploration
 - Transfer Learning
 - Explainability
- We can use surrogate measures to aid option discovery,
 - Bottlenecks/Access States(eg: Doorways)
 - Diverse Density
 - Graph Partitions
 - Betweenness
 - Frequency of Changes
 - Bisimulation Metrics


Finding Bottlenecks

- MDP can be segmented and modeled as a graph.
(Graph: nodes -> states; edge -> action)
- Find components of graph which are weakly connected. (Graph partitioning)
- States where weak connections happen are ‘bottleneck’ states



Diverse Density


- Assume we've obtained a bunch of successful & unsuccessful trajectories. (Experience Replay)
- Find states that appear frequently on successful and rarely on unsuccessful trajectories.
- “The agent needs to get through these states to reach the goal”.



RED: unsuccessful trajectories
GREEN: successful trajectories

Betweenness Centrality

- Pick (all possible) pair of nodes on the graph and calculate the shortest path.
- A node has high betweenness if many shortest-paths pass through it.
- States with high betweenness centrality can be considered bottlenecks.



Small World Options

- Random options inserted with the expected length of these options following a certain probability distribution.
- Exploration time can be cut down significantly.
- Makes best use of data in comparison.